

# PROGRAMMING ASSIGNMENT #2

T81-559: Applications of Deep Neural Networks, Washington University September 16, 2017

Listing 1 shows a sample submission skeleton that you can use as a starting point for this assignment.

Listing 1: Sample Submission Skeleton

```
1
2
3 path = "./data/"
4
5 # Helpful functions
6
7 # Encode text values to dummy variables(i.e. [1,0,0],[0,1,0],[0,0,1] for ↵
   red,green,blue)
8 def encode_text_dummy(df, name):
9     dummies = pd.get_dummies(df[name])
10    for x in dummies.columns:
11        dummy_name = "{}-{}".format(name, x)
12        df[dummy_name] = dummies[x]
13    df.drop(name, axis=1, inplace=True)
14
15
16 # Encode text values to a single dummy variable. The new columns (which ↵
   do not replace the old) will have a 1
17 # at every location where the original column (name) matches each of the ↵
   target_values. One column is added for
18 # each target value.
19 def encode_text_single_dummy(df, name, target_values):
20     for tv in target_values:
21         l = list(df[name].astype(str))
22         l = [1 if str(x) == str(tv) else 0 for x in l]
23         name2 = "{}-{}".format(name, tv)
24         df[name2] = l
25
26
27 # Encode text values to indexes(i.e. [1],[2],[3] for red,green,blue).
28 def encode_text_index(df, name):
29     le = preprocessing.LabelEncoder()
30     df[name] = le.fit_transform(df[name])
31     return le.classes_
32
33
```

```

34 # Encode a numeric column as zscores
35 def encode_numeric_zscore(df, name, mean=None, sd=None):
36     if mean is None:
37         mean = df[name].mean()
38
39     if sd is None:
40         sd = df[name].std()
41
42     df[name] = (df[name] - mean) / sd
43
44
45 # Convert all missing values in the specified column to the median
46 def missing_median(df, name):
47     med = df[name].median()
48     df[name] = df[name].fillna(med)
49
50
51 # Convert all missing values in the specified column to the default
52 def missing_default(df, name, default_value):
53     df[name] = df[name].fillna(default_value)
54
55
56 # Convert a Pandas dataframe to the x,y inputs that TensorFlow needs
57 def to_xy(df, target):
58     result = []
59     for x in df.columns:
60         if x != target:
61             result.append(x)
62     # find out the type of the target column. Is it really this hard? :(
63     target_type = df[target].dtypes
64     target_type = target_type[0] if hasattr(target_type, '__iter__') else ←
        target_type
65     # Encode to int for classification, float otherwise. TensorFlow likes ←
        32 bits.
66     if target_type in (np.int64, np.int32):
67         # Classification
68         dummies = pd.get_dummies(df[target])
69         return df.as_matrix(result).astype(np.float32), dummies.as_matrix(←
            ().astype(np.float32)
70     else:
71         # Regression
72         return df.as_matrix(result).astype(np.float32), df.as_matrix([←
            target]).astype(np.float32)
73
74 # Nicely formatted time string
75 def hms_string(sec_elapsed):
76     h = int(sec_elapsed / (60 * 60))

```

```

77     m = int((sec_elapsed % (60 * 60)) / 60)
78     s = sec_elapsed % 60
79     return "{:}:{:>02}:{:>05.2f}".format(h, m, s)
80
81
82 # Regression chart.
83 def chart_regression(pred,y,sort=True):
84     t = pd.DataFrame({'pred' : pred, 'y' : y.flatten()})
85     if sort:
86         t.sort_values(by=['y'],inplace=True)
87         a = plt.plot(t['y'].tolist(),label='expected')
88         b = plt.plot(t['pred'].tolist(),label='prediction')
89         plt.ylabel('output')
90         plt.legend()
91         plt.show()
92
93 # Remove all rows where the specified column is +/- sd standard deviations
94 def remove_outliers(df, name, sd):
95     drop_rows = df.index[(np.abs(df[name] - df[name].mean()) >= (sd * df[←
        name].std()))]
96     df.drop(drop_rows, axis=0, inplace=True)
97
98
99 # Encode a column to a range between normalized_low and normalized_high.
100 def encode_numeric_range(df, name, normalized_low=-1, normalized_high=1,
101     data_low=None, data_high=None):
102     if data_low is None:
103         data_low = min(df[name])
104         data_high = max(df[name])
105
106     df[name] = ((df[name] - data_low) / (data_high - data_low)) \
107         * (normalized_high - normalized_low) + normalized_low
108
109 # Solution
110
111 def encode_toy_dataset(filename):
112     df = pd.read_csv(filename, na_values=['NA', '?'])
113     encode_numeric_zscore(df, 'length')
114     encode_numeric_zscore(df, 'width')
115     encode_numeric_zscore(df, 'height')
116     encode_text_dummy(df, 'metal')
117     encode_text_dummy(df, 'shape')
118     return df
119
120 # Encode the toy dataset
121 def question1():
122     print()

```

```

123     print("***Question 1***")
124
125     path = "./data/"
126
127     filename_read = os.path.join(path, "toy1.csv")
128     filename_write = os.path.join(path, "submit-jheaton-prog2q1.csv")
129     df = encode_toy_dataset(filename_read) # You just have to implement ↔
        encode_toy_dataset above
130     df.to_csv(filename_write, index=False)
131     print("Wrote {} lines.".format(len(df)))
132
133
134 # Model the toy dataset, no cross validation
135 def question2():
136     print()
137     print("***Question 2***")
138
139 def question3():
140     print()
141     print("***Question 3***")
142
143     # Z-Score encode these using the mean/sd from the dataset (you got ↔
        this in question 2)
144     testDF = pd.DataFrame([
145         {'length':1, 'width':2, 'height': 3},
146         {'length':3, 'width':2, 'height': 5},
147         {'length':4, 'width':1, 'height': 3}
148     ])
149
150
151 def question4():
152     print()
153     print("***Question 4***")
154
155
156 def question5():
157     print()
158     print("***Question 5***")
159
160
161 question1()
162 question2()
163 question3()
164 question4()
165 question5()

```

---

Listing 2 shows what the output from this assignment would look like. Your numbers might differ from mine slightly. Every question, except 2, also generates an output CSV file. For your submission please include your Jupyter notebook and any generated CSV files that the questions specified. Name your output CSV files something such as **submit-jheaton-prog2q1.csv**. Submit a ZIP file that contains your Jupyter notebook and 4 CSV files to Blackboard. This will be 5 files total.

#### Listing 2: Expected Output

---

```
1  ***Question 1***
2  Wrote 10001 lines.
3
4  ***Question 2***
5  Epoch 00144: early stopping
6  Final score (RMSE): 75.46247100830078
7
8  ***Question 3***
9  length: (5.5258474152584744, 2.8609014041584113)
10 width: (5.5340465953404658, 2.8598366585224158)
11 height: (5.5337466253374661, 2.8719829476156122)
12      height    length    width
13 0 -0.882205 -1.581907 -1.235659
14 1 -0.185856 -0.882861 -1.235659
15 2 -0.882205 -0.533338 -1.585338
16
17 ***Question 4***
18 Fold #1
19 Epoch 00060: early stopping
20 Fold score (RMSE): 0.21216803789138794
21 Fold #2
22 Epoch 00061: early stopping
23 Fold score (RMSE): 0.14340682327747345
24 Fold #3
25 Epoch 00028: early stopping
26 Fold score (RMSE): 0.3336745500564575
27 Fold #4
28 Epoch 00058: early stopping
29 Fold score (RMSE): 0.2133668214082718
30 Fold #5
31 Epoch 00077: early stopping
32 Fold score (RMSE): 0.1796143352985382
33 Final, out of sample score (RMSE): 0.22570167481899261
34
35 ***Question 5***
36 Fold #1
37 Epoch 00182: early stopping
```

```
38 Fold score: 0.3625
39 Fold #2
40 Epoch 00425: early stopping
41 Fold score: 0.9875
42 Fold #3
43 Epoch 00169: early stopping
44 Fold score: 0.975
45 Fold #4
46 Epoch 00111: early stopping
47 Fold score: 0.8987341772151899
48 Fold #5
49 Epoch 00203: early stopping
50 Fold score: 0.8227848101265823
51 Final, out of sample score: 0.8090452261306532
```

---

## Question 1

Use the dataset found here for this question: [\[click for toy dataset\]](#).

Encode the **toy1.csv** dataset. Generate dummy variables for the shape and metal. Encode height, width and length as z-scores. Include, but do not encode the weight. If this encoding is performed in a function, named **encode\_toy\_dataset**, you will have an easier time reusing the code from question 1 in question 2.

Write the output to a CSV file that you will submit with this assignment. The CSV file will look similar to Listing 3.

Listing 3: Question 2 Output Sample

---

```
1 height,length,width,weight,metal-bronze,metal-gold,metal-platinum,metal-↵
  silver,metal-tin,shape-box,shape-cylinder,shape-sphere
2 -0.18585564084337075, -0.18381430131795315, -0.1866234261937586, ↵
  729.63,1.0,0.0,0.0,0.0,0.0,0.0,1.0,0.0
3 -0.5340303145851667, -0.18381430131795315, 0.16305509393694917, ↵
  2530.8,0.0,0.0,1.0,0.0,0.0,1.0,0.0,0.0
4 -1.2303796620687586, -1.2323841890415879, -1.5853375067165896, ↵
  58.37,1.0,0.0,0.0,0.0,0.0,0.0,1.0,0.0
5 -0.8822049883269626, -0.8828608931337095, -0.8859804664551741, ↵
  74.15,0.0,0.0,0.0,1.0,0.0,0.0,0.0,1.0
6 ...
```

---

## Question 2

Use the dataset found here for this question: [\[click for toy dataset\]](#).

Use the encoded dataset from question 1 and train a neural network to predict weight. Use 25% of the data as validation and 75% as training, make sure you shuffle the data. Report the RMSE error for the validation set. No CSV file need be generated for this question.

### Question 3

Use the dataset found here for this question: [\[click for toy dataset\]](#).

Using the **toy1.csv** dataset calculate and report the mean and standard deviation for height, width and length. Calculate the z-scores for the dataframe given by Listing 4. Make sure that you use the mean and standard deviations you reported for this question. Write the results to a CSV file.

Listing 4: Question 3 Input Data

```
1 testDF = pd.DataFrame([
2     {'length':1, 'width':2, 'height': 3},
3     {'length':3, 'width':2, 'height': 5},
4     {'length':4, 'width':1, 'height': 3}
5 ])
6 ...
```

Your resulting CSV file should look almost exactly like Listing 5.

Listing 5: Question 3 Output Sample

```
1 height,length,width
2 -0.8822049883269626,-1.5819074849494659,-1.2356589865858818
3 -0.18585564084337075,-0.8828608931337095,-1.2356589865858818
4 -0.8822049883269626,-0.5333375972258314,-1.5853375067165896
```

### Question 4

Use the dataset found here for this question: [\[click for iris dataset\]](#).

Usually the **iris.csv** dataset is used to classify the species. Not this time! Use the fields species, sepal-l, sepal-w, and petal-l to predict petal-w. Use a 5-fold cross validation and report ONLY out-of-sample predictions to a CSV file. Make sure to shuffle the data. Your generated CSV file should look similar to Listing 6. Encode each of the inputs in a way that makes sense (e.g. dummies, z-scores).

Listing 6: Question 4 Output Sample

```
1 sepal_l,sepal_w,petal_l,petal_w,species-Iris-setosa,species-Iris-↵
   versicolor,species-Iris-virginica,0,0
```

```

2 0.30995914214417364, -0.5903951331558184, 0.5336208818725668, ↵
    1.2,0.0,1.0,0.0,1.2,1.444551944732666
3 -0.1730940663922016, 1.7038864723719687, -1.1658086782311483, ↵
    0.3,1.0,0.0,0.0,0.3,0.\
4 ...

```

---

## Question 5

Use the dataset found here for this question: [click for auto mpg dataset].

Usually the **auto-mpg.csv** dataset is used to regress the mpg. Not this time! Use the fields to predict how many cylinders the car has. Treat this as a classification problem, where there is a class for each number of cylinders. Use a 5-fold cross validation and report ONLY out-of-sample predictions to a CSV file. Make sure to shuffle the data. Your generated CSV file should look similar to Listing 7. Encode each of the inputs in a way that makes sense (e.g. dummies, z-scores). Report the final out of sample accuracy score.

Listing 7: Question 4 Output Sample

```

1 mpg,cylinders,displacement,horsepower,weight,acceleration,year,origin,name↵
    ,ideal,predict
2 -0.7055506566787514, 8, 1.0892327311042995, 0.6722714619460141, ↵
    0.6300768256149949, -1.2938698102195594, 70, -0.7142457922976494,↵
    chevrolet chevelle malibu,8,8
3 -1.0893794720944747, 8, 1.5016242793620063, 1.5879594901955474, ↵
    0.8532590135498572, -1.4751810504376373, 70, -0.7142457922976494,buick↵
    skylark 320,8,8
4 -0.7055506566787514, 8, 1.1947282434492943, 1.19552176380289, ↵
    0.5497784722839334, -1.6564922906557151, 70, -0.7142457922976494,↵
    plymouth satellite,8,8
5 ...

```

---